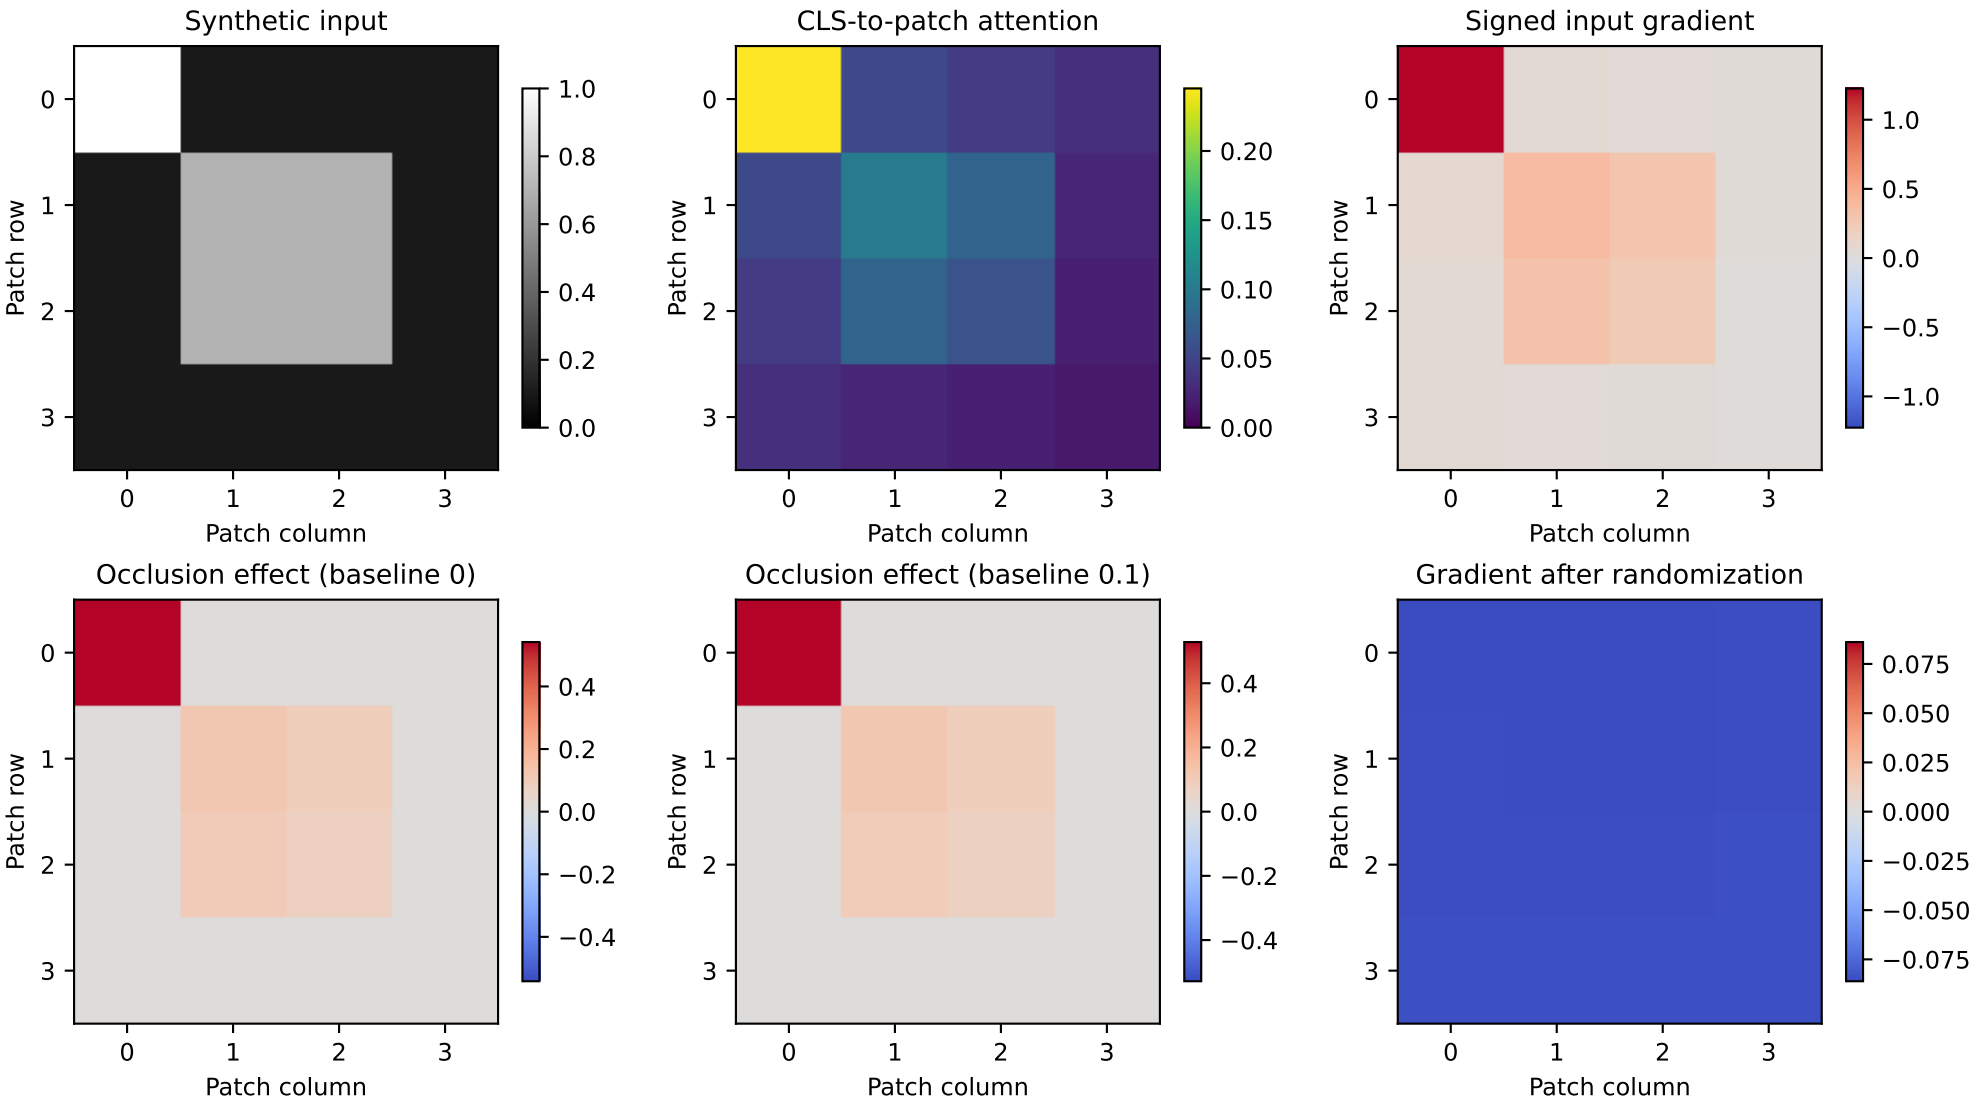
Same input, different questions: a hand-specified attention sublayer
Synthetic input
CLS-to-patch attention
Signed input gradient
Occlusion effect (baseline 0)
Occlusion effect (baseline 0.1)
Gradient after randomization
Patch row
Patch column
0
1
2
3
1.0
0.8
0.6
0.4
0.2
0.0
0.20
0.15
0.10
0.05
0.00
0.5
−0.5
−1.0
−0.2
−0.4
0.075
0.050
0.025
0.000
−0.025
−0.050
−0.075